

Databases

Lecture 6

Functional Dependencies, Normal Forms

Relational Algebra

Functional Dependencies, Normal Forms

- a dependency (simple, multi-valued) in a relation can be eliminated via decompositions (the original relation is decomposed into a collection of new relations)
- nevertheless, there are relations without such dependencies that can still contain redundant information, which can be a source of errors in the database

->

Example 13. Consider the relation FaPrCo [FacultyMember, Program, Course], storing the programs and courses for different faculty members; this relation has no functional dependencies; its key is {FacultyMember, Program, Course}

- consider the following data in the relation:

Fa	Pr	Co
F1	P1	C2
F1	P2	C1
F2	P1	C1
F1	P1	C1

- some associations appear in multiple records (redundant data):
 - faculty member F1 is teaching in program P1
 - faculty member F1 is teaching course C1
 - course C1 is taught in program P1

Fa	Pr	Co
F1	P1	C2
F1	P2	C1
F2	P1	C1
F1	P1	C1

- if some values in the relation are changed, e.g., "F1 will teach course C3 instead of course C1", several updates should be carried out, without knowing in how many records
- the same is true for the following changes: "in P1, course C3 should replace course C1", "F1 is switching from program P1 to P3"

- the previous relation cannot be decomposed into 2 relations (via projection), because new data would be introduced through the join
- this claim can be justified by considering the three possible projections on two attributes:

FaPr	Fa	Pr
	F1	P1
	F1	P2
	F2	P1

FaCo	Fa	Co
	F1	C2
	F1	C1
	F2	C1

PrCo	Pr	Co
	P1	C2
	P2	C1
	P1	C1

- when evaluating $\text{FaPr} * \text{PrCo}$, the following data is obtained:

$R' = \text{FaPr} * \text{PrCo}$	Fa	Pr	Co
	F1	P1	C2
	F1	P1	C1
	F1	P2	C1
	F2	P1	C2
	F2	P1	C1

- this result set contains an extra tuple, which didn't exist in the original relation
- this is also true for the other join combinations: $\text{FaPr} * \text{FaCo}$ and $\text{PrCo} * \text{FaCo}$

- when evaluating $R' * FaCo$ (i.e., $FaPr * PrCo * FaCo$), the original relation $FaPrCo$ is obtained
- conclusion: $FaPrCo$ cannot be decomposed into 2 projections, but it can be decomposed into 3 projections, i.e., $FaPrCo$ is *3-decomposable*:

$$FaPrCo = FaPr * PrCo * FaCo, \text{ or } FaPrCo = * (FaPr, PrCo, FaCo)$$

- this conclusion ($FaPrCo$ is 3-decomposable) is true for the data in the relation
- the 3-decomposability can be specified as a restriction that must be met by all the legal instances of the relation:

* if $(F1, P1) \in FaPr$ and $(F1, C1) \in FaCo$ and $(P1, C1) \in PrCo$ then $(F1, P1, C1) \in FaPrCo$

- this restriction can be expressed on $FaPrCo$:

* if $(F1, P1, C2) \in FaPrCo$ and $(F1, P2, C1) \in FaPrCo$ and $(F2, P1, C1) \in FaPrCo$ then $(F1, P1, C1) \in FaPrCo$

- consider the following data in the relation

Fa	Pr	Co
F1	P1	C2
F1	P2	C1

- if the previous restriction is specified, then, if (F2, P1, C1) is added to the relation, (F1, P1, C1) must be also added:

Fa	Pr	Co
F1	P1	C2
F1	P2	C1
F2	P1	C1
F1	P1	C1

- if (F1, P1, C1) is removed from the relation, other data must be removed as well, at least (F2, P1, C1), for the restriction to be satisfied

Definition. Let $R[A]$ be a relation and $R_i[\alpha_i]$, $i=1,\dots,m$, the projections of R on α_i . R satisfies the join dependency $* \{\alpha_1, \dots, \alpha_m\}$ if $R = R_1 * \dots * R_m$.

- FaPrCo has a join dependency because $\text{FaPrCo} = \text{FaPr} * \text{PrCo} * \text{FaCo}$
- JD $* \{\alpha_1, \dots, \alpha_m\}$ is trivial if at least one of α_i is the set of all attributes in R .
- JD $* \{\alpha_1, \dots, \alpha_m\}$ is implied by the candidate keys if each α_i is a superkey in R .

Definition. Relation R is in 5NF if every non-trivial JD is implied by the candidate keys in R .

- $R[A]$ - a relation
- F - a set of functional dependencies
- α – a subset of attributes

- problems

I. compute the closure of F : F^+

II. compute the closure of a set of attributes under a set of functional dependencies, e.g., the closure of α under F : α^+

III. compute the minimal cover for a set of dependencies

- $R[A]$ - a relation
- F - a set of functional dependencies
- problems

I. compute the closure of F : F^+

- the set F^+ contains all the functional dependencies implied by F
- F implies a functional dependency f if f holds on every relation that satisfies F
- the following 3 rules can be repeatedly applied to compute F^+ (Armstrong's axioms):
 - $\alpha, \beta, \gamma \subset A$
 1. reflexivity: if $\beta \subseteq \alpha$, then $\alpha \rightarrow \beta$
 2. augmentation: if $\alpha \rightarrow \beta$, then $\alpha\gamma \rightarrow \beta\gamma$
 3. transitivity: if $\alpha \rightarrow \beta$ and $\beta \rightarrow \gamma$, then $\alpha \rightarrow \gamma$
- these rules are complete (they compute the closure) and sound (no erroneous functional dependencies can be derived)

- $R[A]$ - a relation
- F - a set of functional dependencies
- problems

I. compute the closure of F : F^+

- the following rules can be derived from Armstrong's axioms:

4. union: if $\alpha \rightarrow \beta$ and $\alpha \rightarrow \gamma$, then $\alpha \rightarrow \beta\gamma$

$$\left. \begin{array}{l} \alpha \rightarrow \beta \Rightarrow \alpha\alpha \rightarrow \alpha\beta \\ \text{augmentation} \\ \alpha \rightarrow \gamma \Rightarrow \alpha\beta \rightarrow \beta\gamma \\ \text{augmentation} \end{array} \right\} \begin{array}{l} \Rightarrow \\ \text{transitivity} \end{array} \alpha \rightarrow \beta\gamma$$

5. decomposition: if $\alpha \rightarrow \beta\gamma$, then $\alpha \rightarrow \beta$ and $\alpha \rightarrow \gamma$

$$\left. \begin{array}{l} \alpha \rightarrow \beta\gamma \\ \beta\gamma \rightarrow \beta \text{ (reflexivity)} \end{array} \right\} \begin{array}{l} \Rightarrow \\ \text{transitivity} \end{array} \alpha \rightarrow \beta \text{ (}\alpha \rightarrow \gamma \text{ can similarly be shown to hold)}$$

- $R[A]$ - a relation
- F - a set of functional dependencies
- problems

I. compute the closure of F : F^+

- the following rules can be derived from Armstrong's axioms:

6. pseudo transitivity: if $\alpha \rightarrow \beta$ and $\beta\gamma \rightarrow \delta$, then $\alpha\gamma \rightarrow \delta$

$$\left. \begin{array}{l} \alpha \rightarrow \beta \\ \beta\gamma \rightarrow \delta \end{array} \right\} \begin{array}{l} \Rightarrow \\ \text{transitivity} \end{array} \alpha\gamma \rightarrow \delta$$

- $\alpha, \beta, \gamma, \delta \subset A$

- $R[A]$ - a relation
- F - a set of functional dependencies
- α – a subset of attributes
- problems

II. compute the closure of a set of attributes under a set of functional dependencies

- determine the closure of α under F , denoted by α^+
- α^+ - the set of attributes that are functionally dependent on attributes in α (under F)

- $R[A]$ - a relation
- F - a set of functional dependencies
- α – a subset of attributes
- problems

II. compute the closure of a set of attributes under a set of functional dependencies

- algorithm:

closure $:= \alpha$;

repeat until there is no change

for every functional dependency $\beta \rightarrow \gamma$ **in** F

if $\beta \subseteq \text{closure}$

then **closure** $:= \text{closure} \cup \gamma$;

- $R[A]$ - a relation
- F - a set of functional dependencies
- problems

III. compute the minimal cover for a set of dependencies

Definition: F, G - two sets of functional dependencies; F and G are equivalent (notation $F \equiv G$) if $F^+ = G^+$.

->

- $R[A]$ - a relation
- F - a set of functional dependencies
- problems

III. compute the minimal cover for a set of dependencies

Definition: F - set of functional dependencies; a minimal cover for F is a set F_M of functional dependencies such that:

1. $F_M \equiv F$
2. the right side of every dependency in F_M has a single attribute;
3. the left side of every dependency in F_M is irreducible (i.e., no attribute can be removed from the determinant of a dependency in F_M without changing F_M 's closure);
4. no dependency f in F_M is redundant (no dependency can be discarded without changing F_M 's closure).

- see lecture problems
 - R - a relation, F - a set of functional dependencies, f - a functional dependency
 - show that f is in F^+
 - R - a relation, F - a set of functional dependencies, α - a subset of the attributes of R
 - compute α^+
 - R - a relation, F - a set of functional dependencies
 - compute F_M - minimal cover for F

Relational Algebra

- query languages in the relational model
 - relational algebra and calculus - formal query languages with a significant influence on SQL
 - relational algebra
 - queries are specified in an operational manner
 - relational calculus
 - queries describe the desired answer, without specifying how it will be computed (declarative)
 - not expected to be Turing complete
 - not intended for complex calculations
 - provide smooth, efficient access to large datasets
 - allow optimizations

- relational algebra
 - used by DBMSs to represent query execution plans
 - a relational algebra query:
 - is built using a collection of operators
 - describes a step-by-step procedure for computing the result set
 - is evaluated on the input relations' instances
 - produces an instance of the output relation
 - every operation returns a relation, hence operators can be composed; the algebra is closed
 - operating on sets of tuples; extension - duplicates are not eliminated

Conditions

- conditions that can be used in several algebraic operators
- similar to the SELECT filter conditions

1. *attribute_name relational_operator value*

- *value* - attribute name, expression

2. *attribute_name IS [NOT] IN single_column_relation*

- a relation with one column can be considered a set
- the condition tests whether a value belongs to a set

3. *relation {IS [NOT] IN | = | <>} relation*

- the relations in the condition should be union-compatible

Conditions

4. *(condition)*

NOT condition

condition₁ AND condition₂

condition₁ OR condition₂,

where *condition*, *condition₁*, *condition₂* are conditions of type 1-4.

Operators in the Algebra

- equivalent SELECT statements can be specified for the relational algebra expressions
- *selection*
 - unary operator
 - notation: $\sigma_C(R)$
 - resulting relation:
 - schema: R's schema
 - tuples: records in R that meet condition C
 - equivalent SELECT statement
 - `SELECT * FROM R WHERE C`

- *projection*
 - unary operator
 - notation: $\Pi_{\alpha}(R)$
 - resulting relation:
 - schema: attributes in α
 - tuples: every record in R is projected on α
 - α can be extended to a set of expressions, specifying the columns of the relation being computed
 - equivalent SELECT statement
 - `SELECT α FROM R`

References

- [Ta13] ȚÂMBULEA, L., Curs Baze de date, Facultatea de Matematică și Informatică, UBB, 2013-2014
- [Ra00] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (2nd Edition), McGraw-Hill, 2000
- [Da03] DATE, C.J., An Introduction to Database Systems (8th Edition), Addison-Wesley, 2003
- [Ga08] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book, Prentice Hall Press, 2008
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007,
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems,
<http://infolab.stanford.edu/~ullman/fcdb.html>