

Databases

Lecture 10


Indexes. Trees

Indexed Sequential Access Method (ISAM)

- insertion
 - find the corresponding leaf page, add the entry
 - if there is no space on the page, add an overflow page
- deletion
 - find the leaf page that contains the entry, remove the entry
 - if an overflow page is emptied, it can be eliminated
- inserts / deletes
 - only leaf pages are affected (static structure)

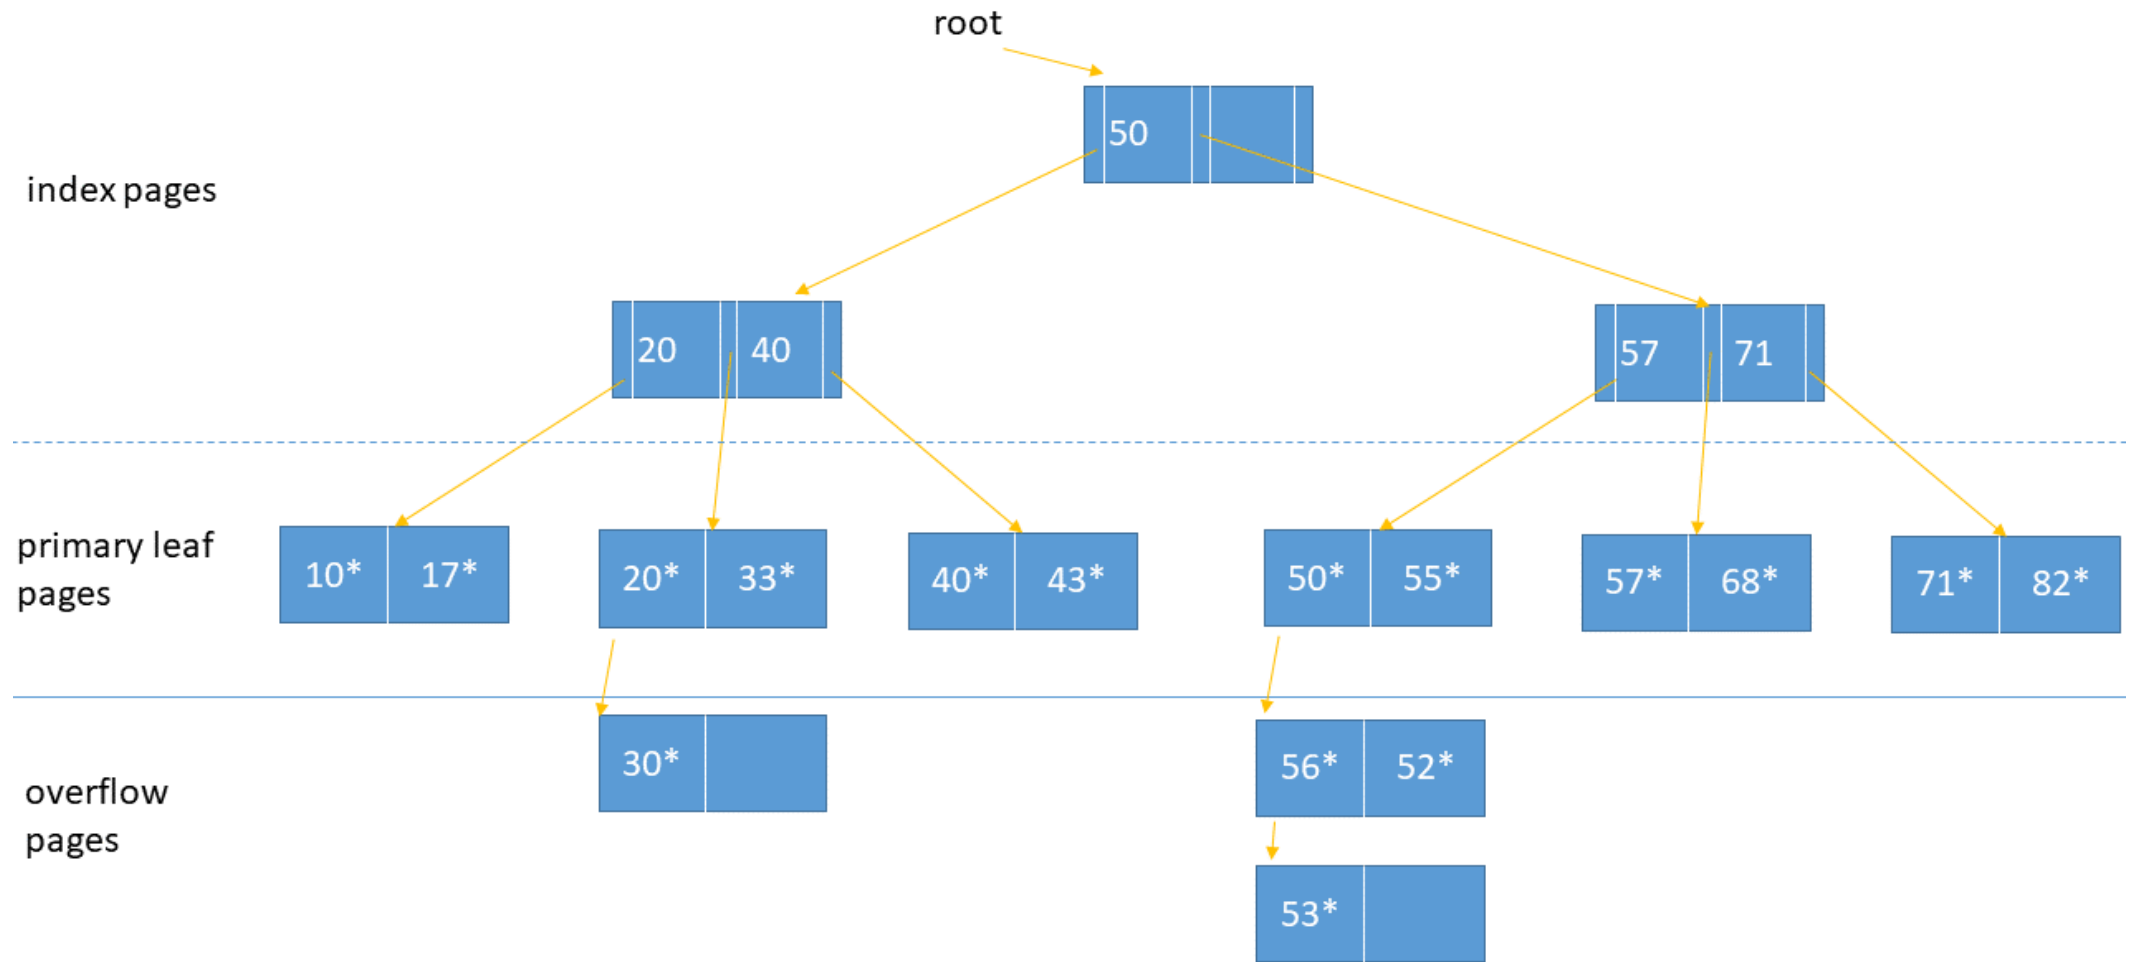
* Example - ISAM tree

- leaf page - 2 entries

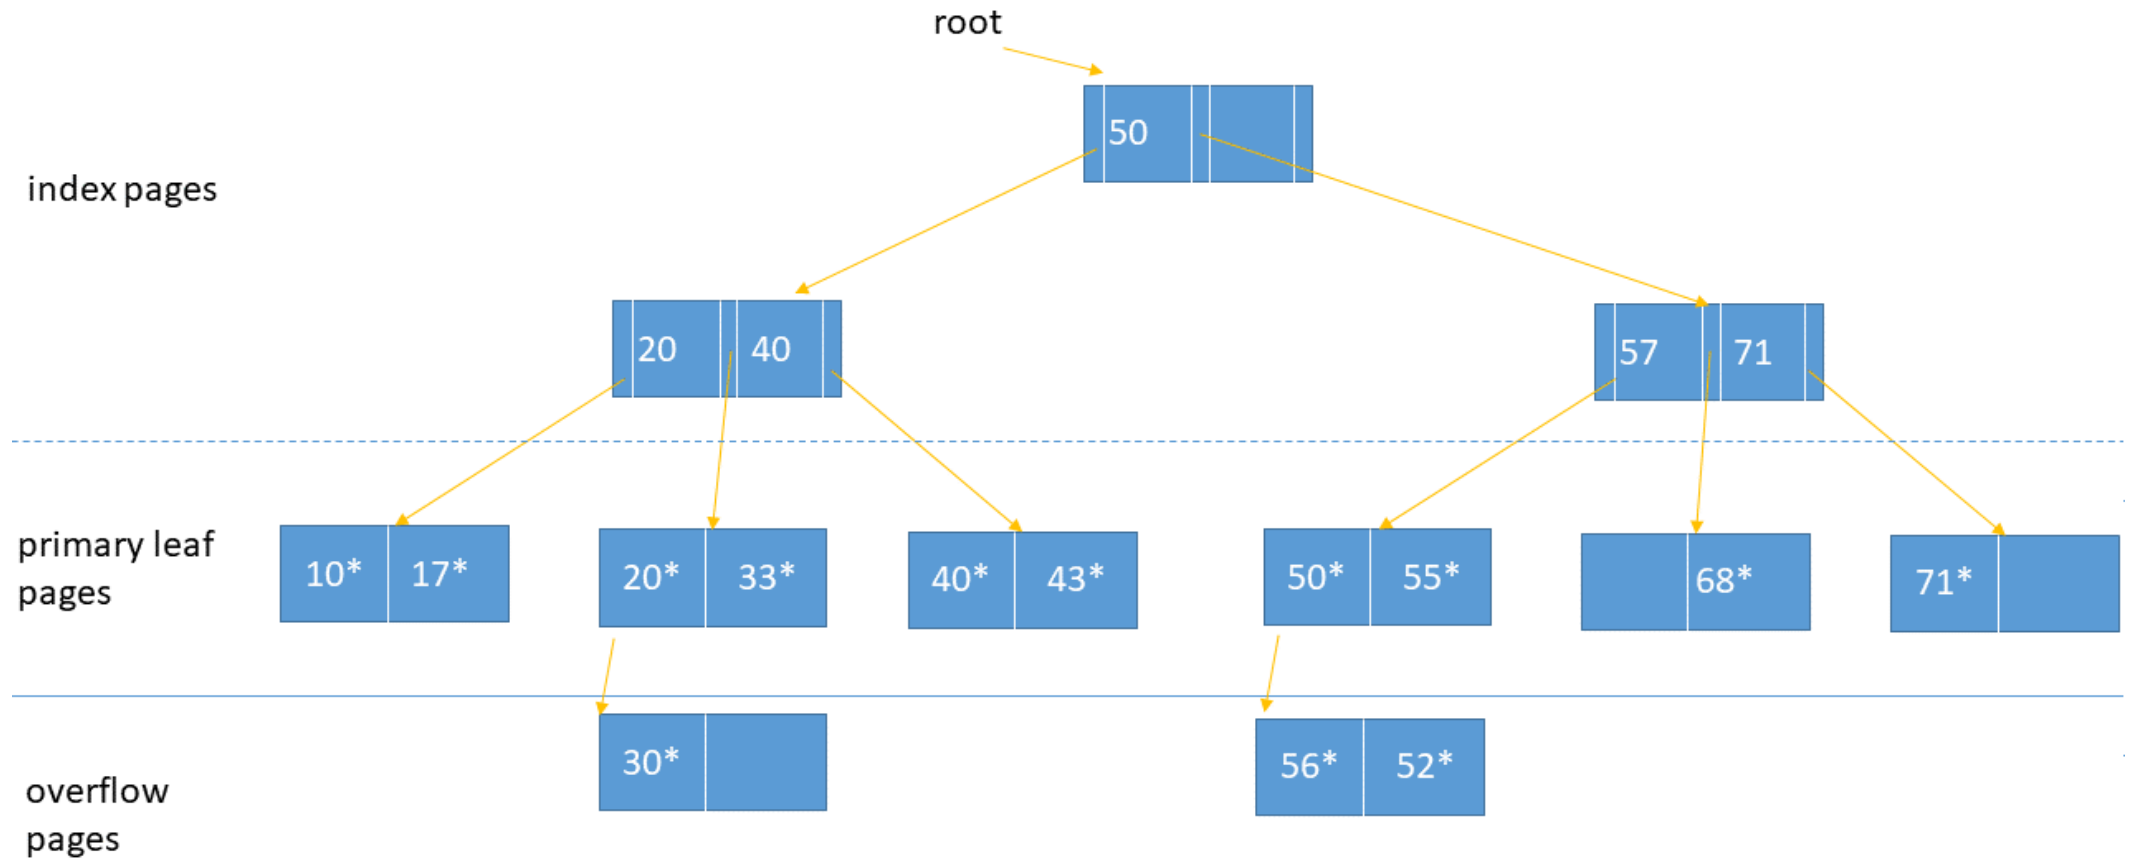


- only key values are shown

- after inserting 30*, 56*, 52*, 53*



- after deleting 53*, 57*, 82*



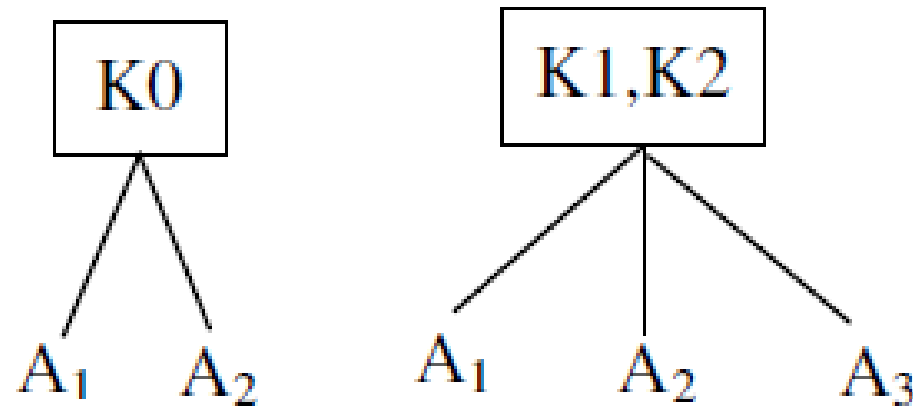
- even if it doesn't appear in the leaves, 57 still appears in an inner node

- benefits and drawbacks
 - inserts / deletes - faster
 - no balancing
 - no I/O operations (changes) on inner nodes
 - better concurrent access, since only leaf pages are modified
 - long overflow chains can develop
 - usually not sorted (to optimize inserts)
 - irregular search time if structure not balanced
 - at file creation - 20% of each page is free for future insertions
 - eliminated through deletes / file reorganization
 - ISAM - suitable when data size and distribution is relatively static


2-3 tree

2-3 tree storing key values (collection of distinct values)

- all the terminal nodes are on the same level
- every node has 1 or 2 key values
 - a non-terminal node with one value K_0 has 2 subtrees: one with values less than K_0 , and one with values greater than K_0
 - a non-terminal node with 2 values K_1 and K_2 , $K_1 < K_2$, has 3 subtrees: one with values less than K_1 , a subtree with values between K_1 and K_2 , and a subtree with values greater than K_2



* Example (key values are letters)




- storing a 2-3 tree
 - 2-3 tree index storing the values of a key
 - tree - key value + address of record (file / DB address of record with corresponding key value)

- 2 options

1. transform 2-3 tree into a binary tree

- nodes with 2 values are transformed (see figure below)
- nodes with 1 value - unchanged



- the structure of a node

K	ADDR	PointerL	PointerR	IND
---	------	----------	----------	-----

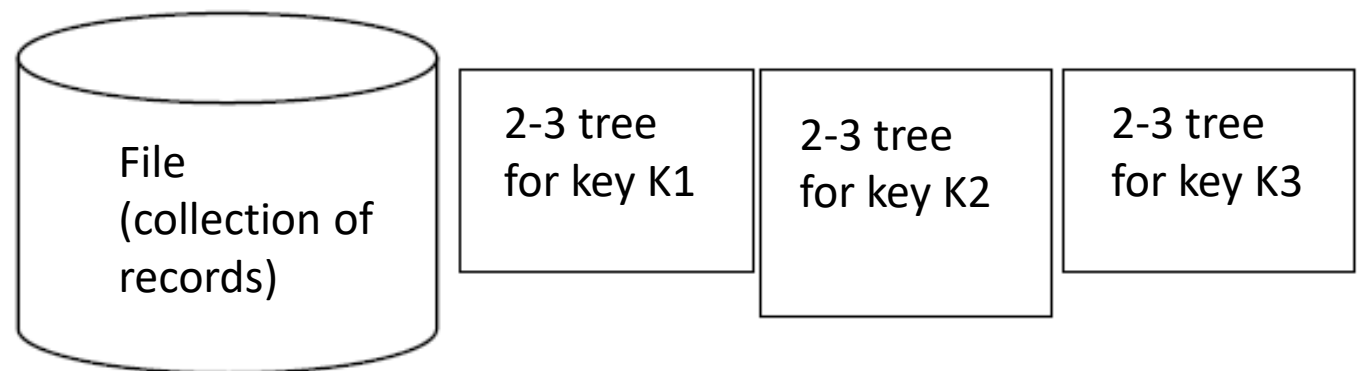
- K - key value
- ADDR - address of the record with the current key value (address in the file)
- PointerL, PointerR - the 2 subtrees' addresses (address in the tree)

- IND - indicator that specifies the type of the link to the right (the 2 possible values can be seen in the previous figure)

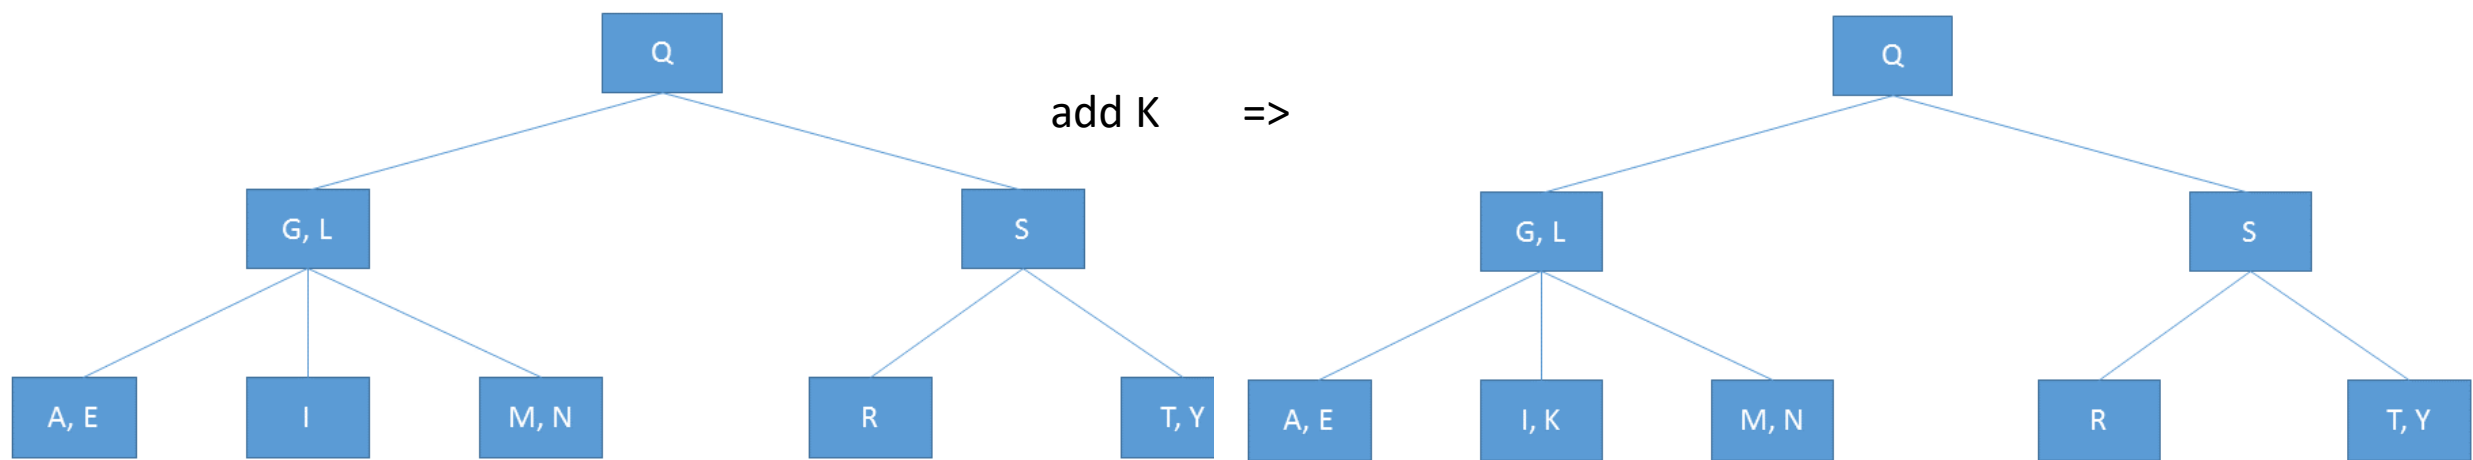
2. the memory area allocated for a node can store 2 values and 3 subtree addresses

NV	K_1	ADDR ₁	K_2	ADDR ₂	Pointer ₁	Pointer ₂	Pointer ₃
----	-------	-------------------	-------	-------------------	----------------------	----------------------	----------------------

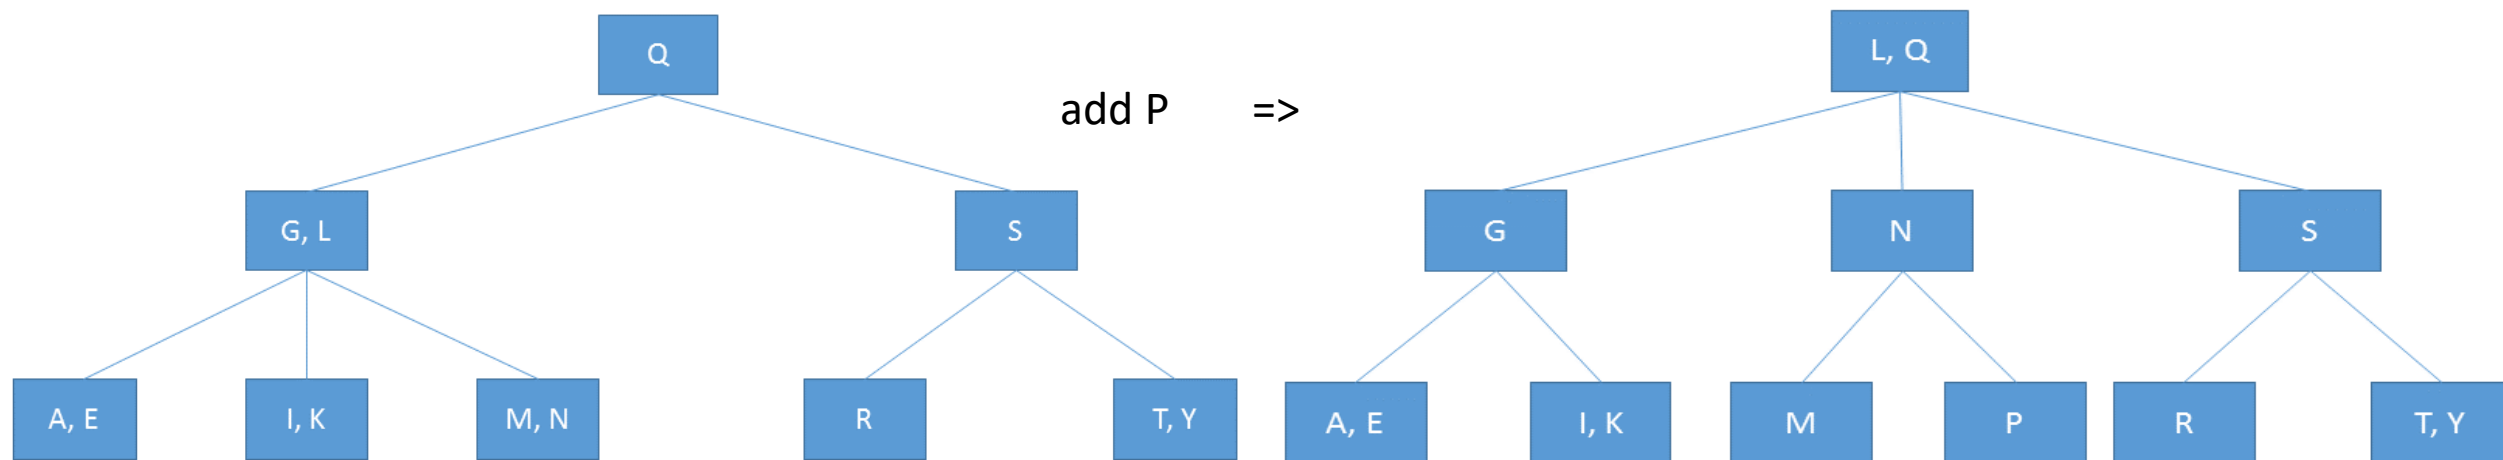
- NV – number of values in the node (1 or 2)
- K_1, K_2 – key values
- ADDR₁, ADDR₂ – the records' addresses (corresponding to K_1 and K_2)
- Pointer₁, Pointer₂, Pointer₃ – the 3 subtrees' addresses
- obs. a file (a relation in a relational DB) can have several associated 2-3 trees (e.g., one tree / key)



- operations in a 2-3 tree
 - searching for a record with key value K_0
 - inserting a record
 - removing a record
 - tree traversal (partial, total)
- add a new value
 - values in the tree must be distinct, i.e., the new value should not exist in the tree
 - perform a test, i.e., search for the value in the tree; if the new value can be added, the search ends in a terminal node
 - if the reached terminal node has 1 value, the new value can be stored in the node



- if the reached terminal node has 2 values, the new value is added to the node, the 3 values are sorted, the node is split into 2 nodes: one node will contain the smallest value, the 2nd node - the largest value, and the middle value is attached to the parent node; the parent is then analyzed in a similar manner

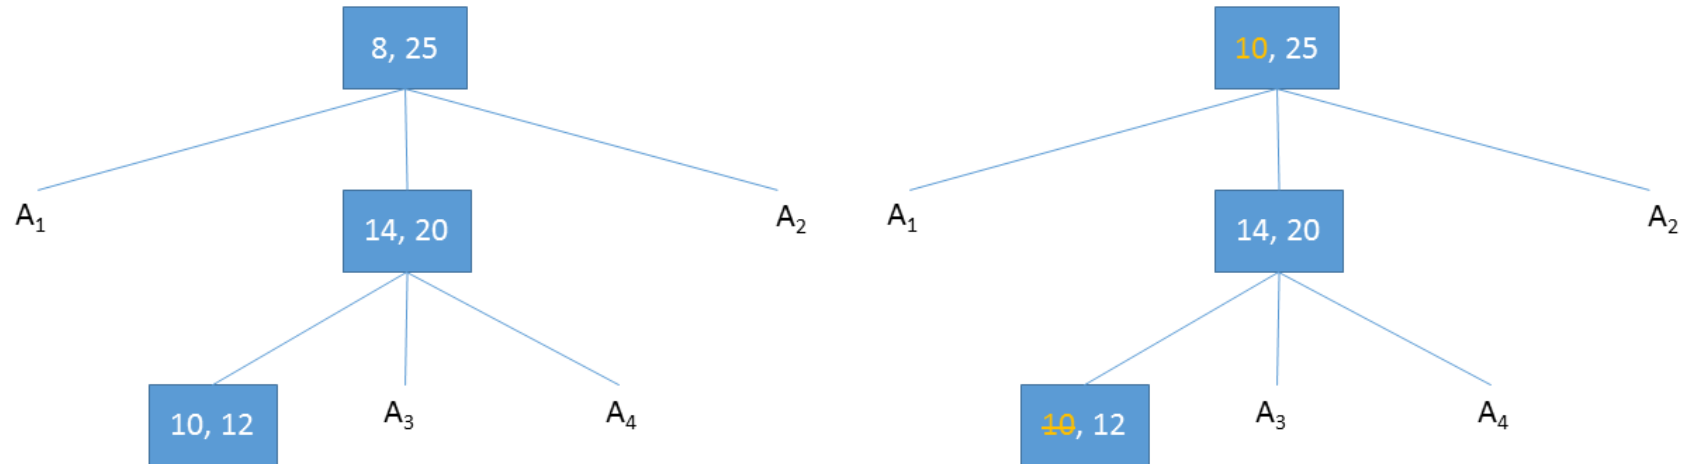


- delete a value K_0

1. search for K_0 ; if K_0 appears in an inner node, change it with a neighbor value K_1 from a terminal node (there is no other value between K_0 and K_1)

- K_1 's previous position (in the terminal node) is eliminated

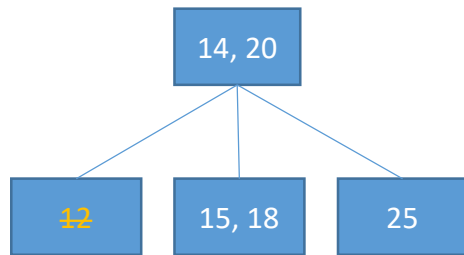
- e.g., remove 8:




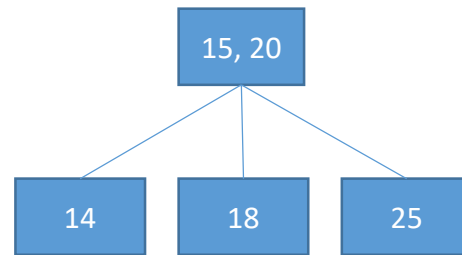
2. perform this step until case a / b occurs

a. if the current node (from which a value is removed) is the root or a node with 1 remaining value, the value is eliminated; the algorithm ends ->

b. if the delete operation empties the current node (it has no values), but 2 values exist in one of the sibling nodes (left / right), 1 of the sibling's values is transferred to the parent, 1 of the parent's values is transferred to the current node; the algorithm ends



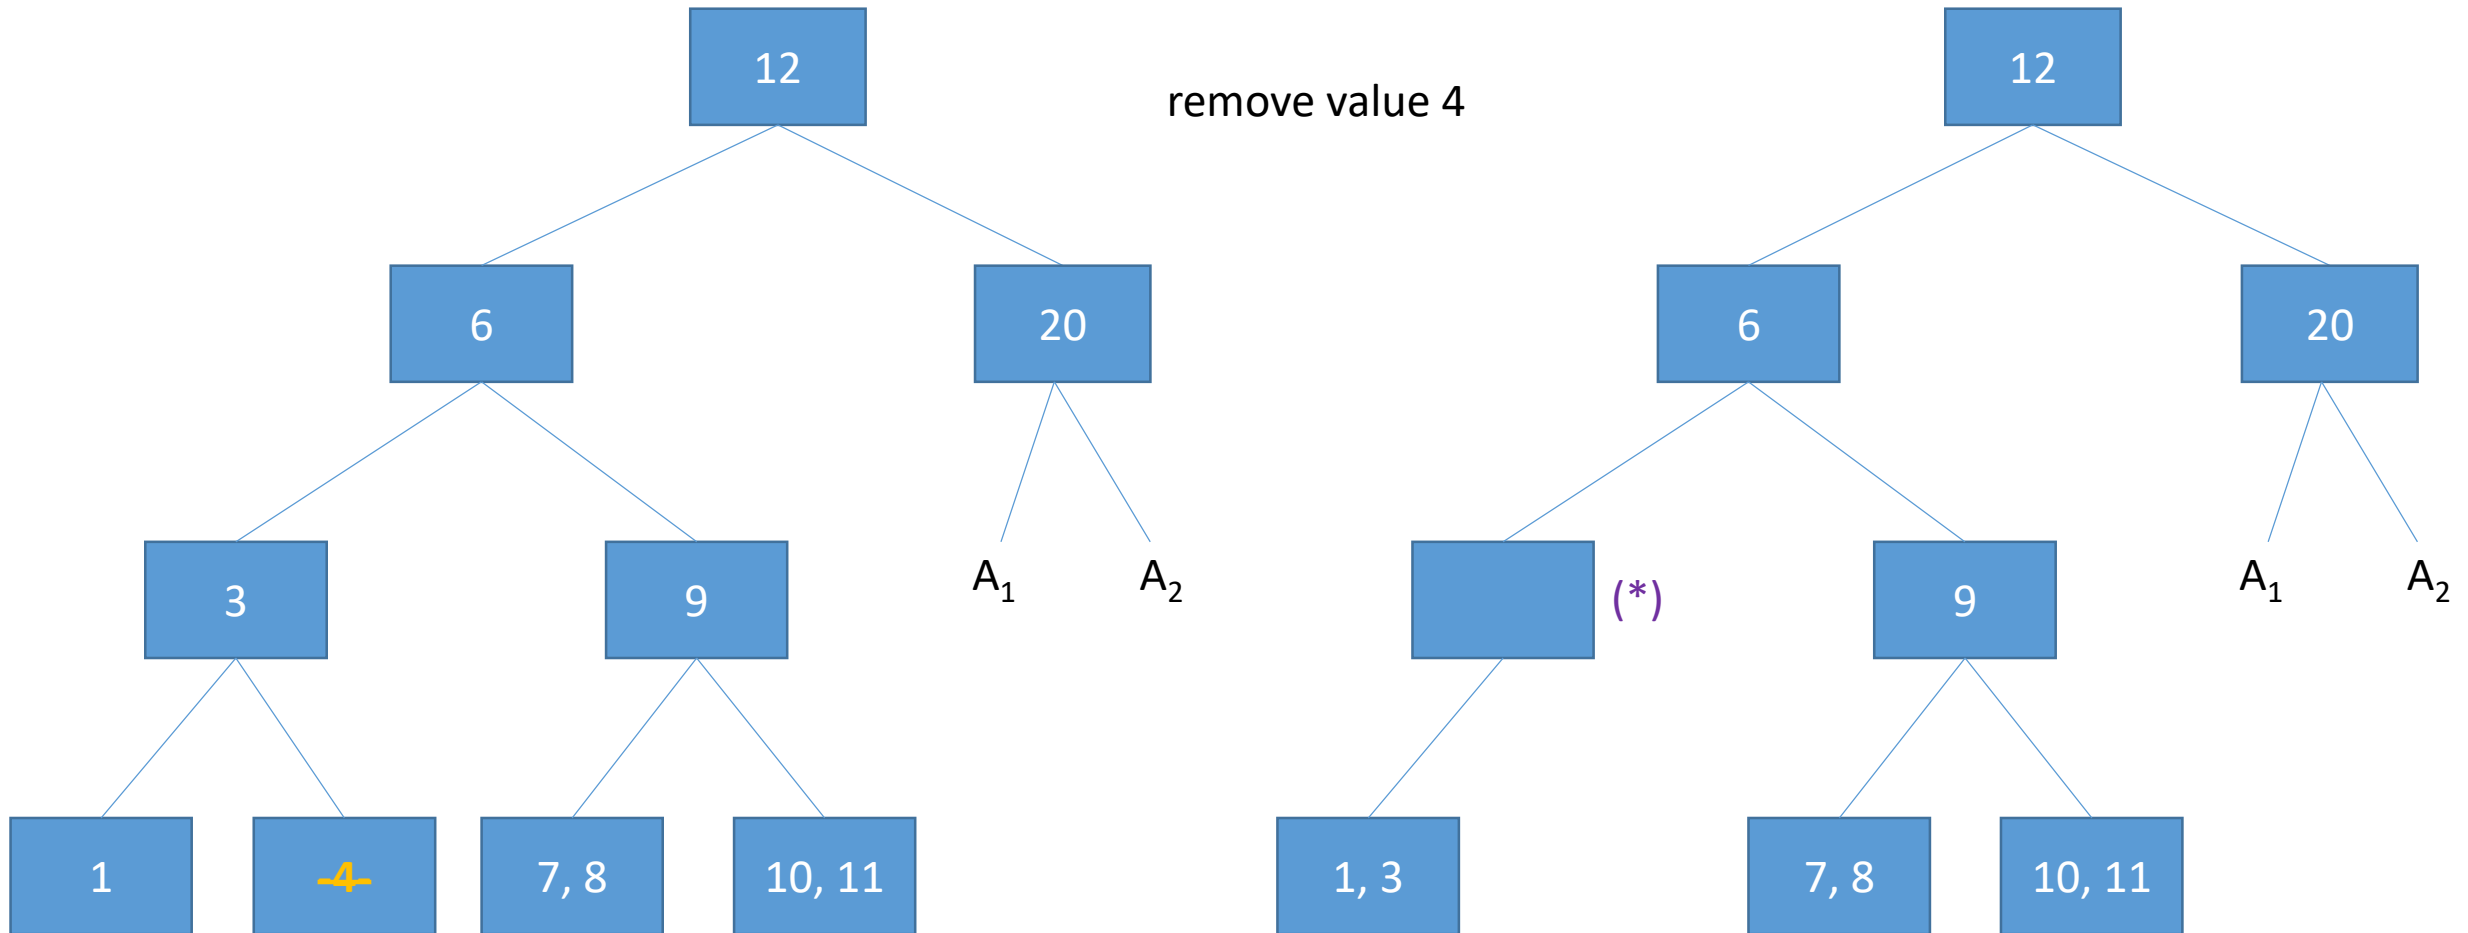
delete value 12

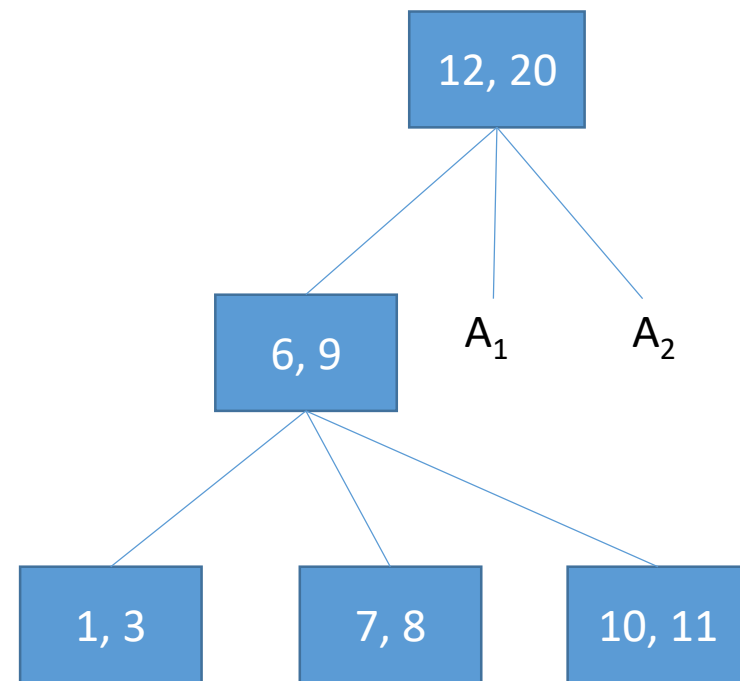
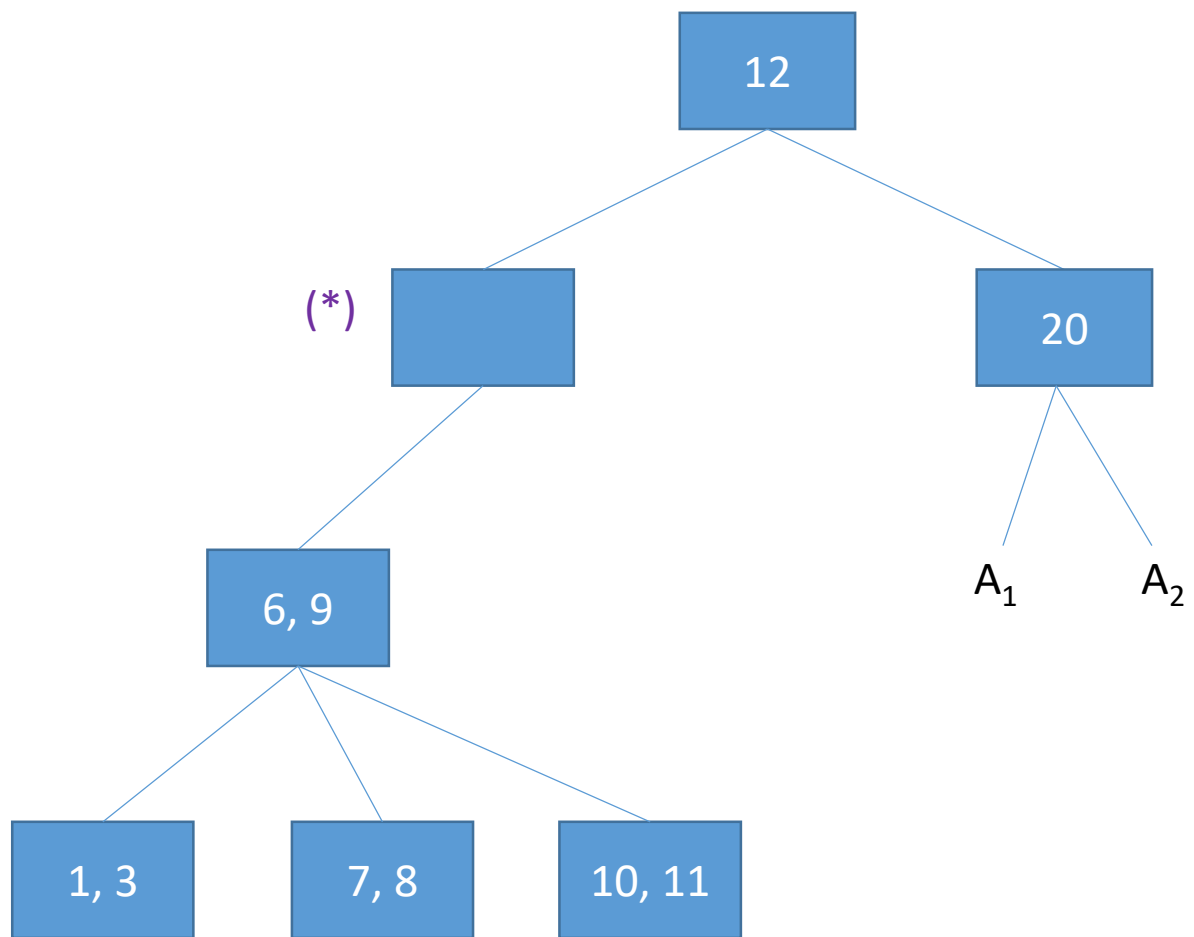


c. if the previous cases do not occur (current node has no values, sibling nodes have 1 value each), then the current node is merged with a sibling and a value from the parent node; case 2 is then analyzed for the parent

- if the root is reached and it has no values, it is eliminated and the current node becomes the root

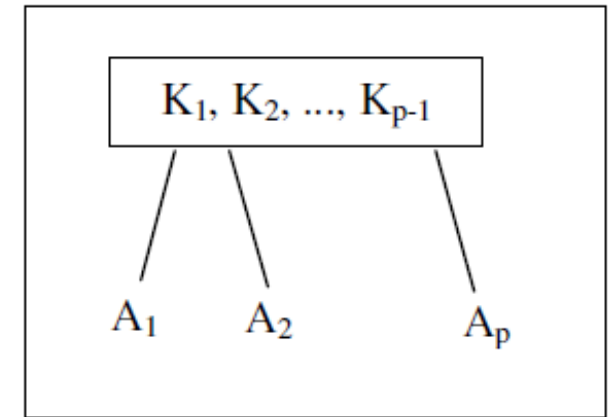
- e.g., case c for the node marked with (*)





B-tree

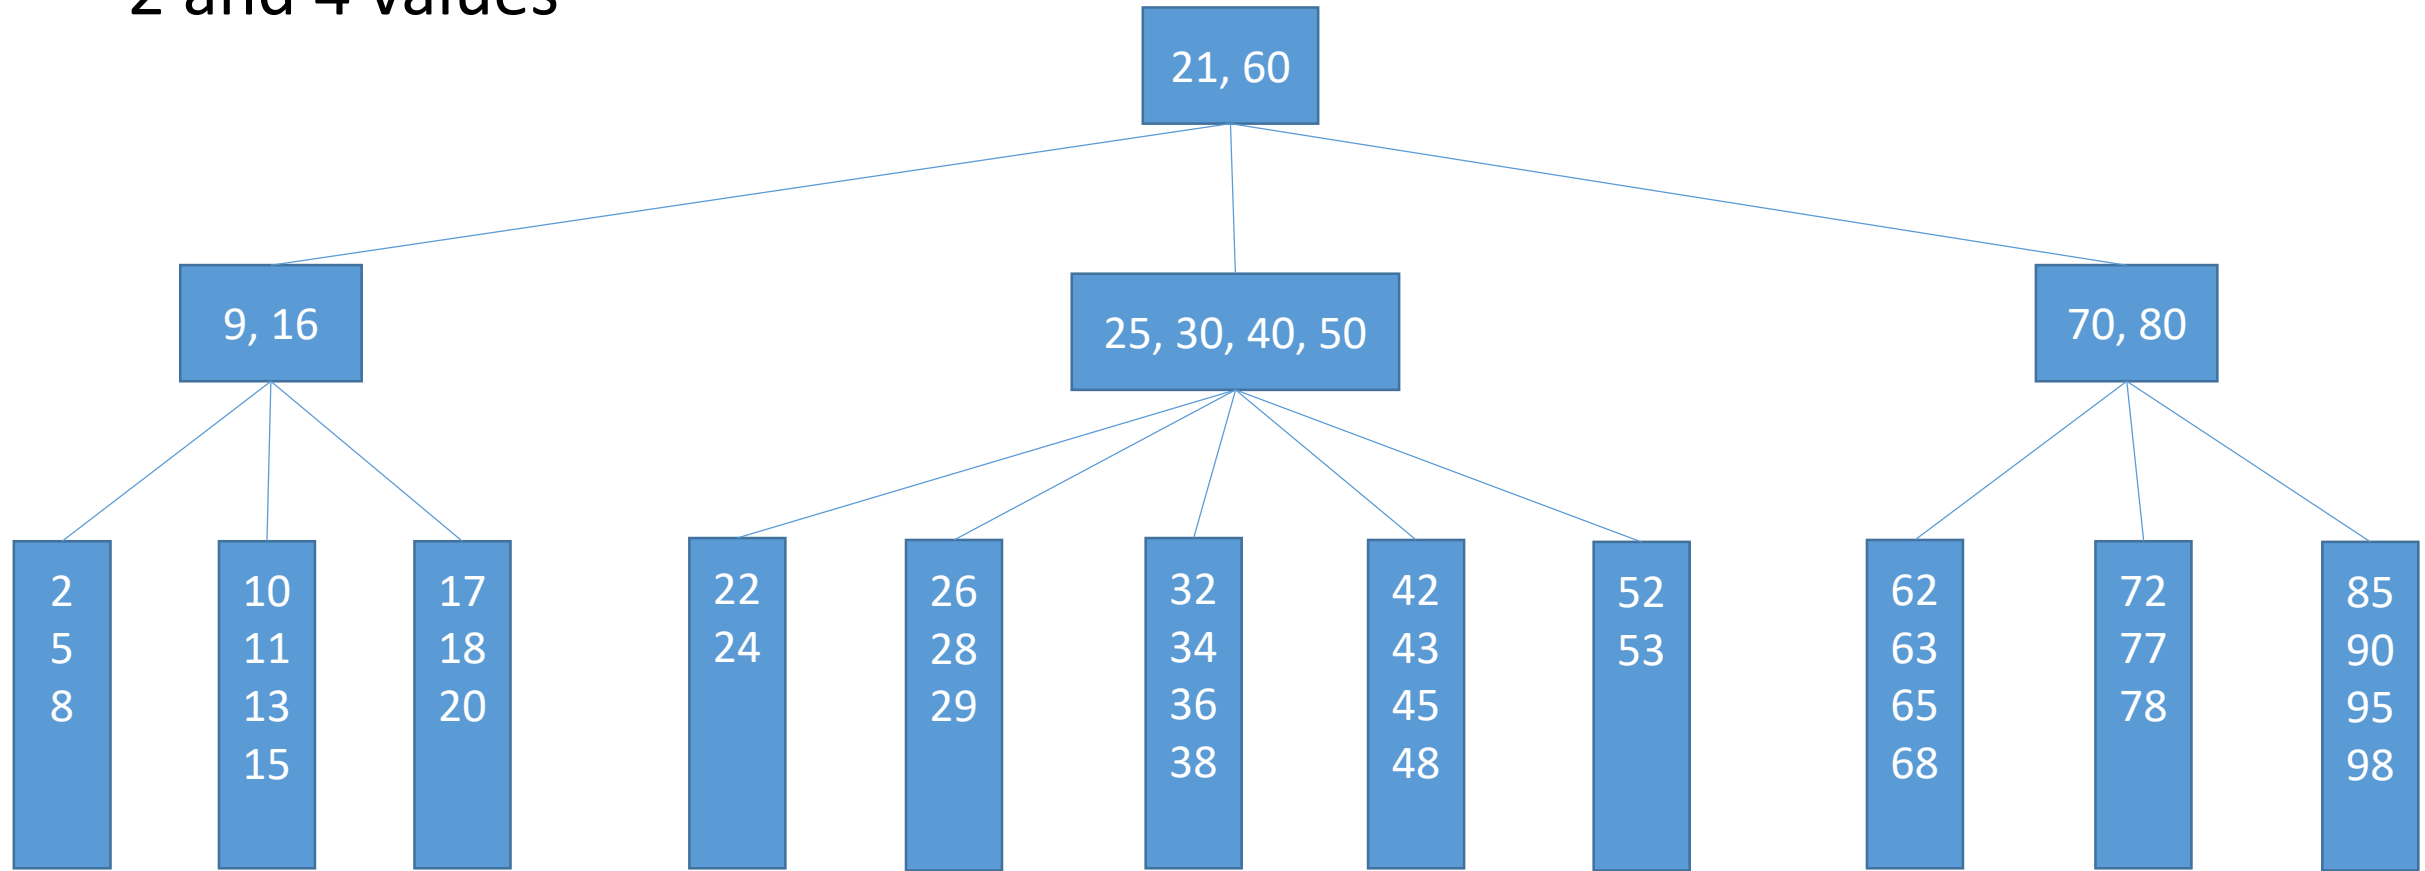
- generalization of 2-3 trees
- B-tree of order m
 1. if the root is not a terminal, it has at least 2 subtrees
 2. all terminal nodes – same level
 3. every non-terminal node – at most m subtrees
 4. a node with p subtrees has p-1 ordered values (ascending order), i.e., $K_1 < K_2 < \dots < K_{p-1}$
 - A_1 : values less than K_1
 - A_i : values between K_{i-1} and K_i , $i=2, \dots, p-1$
 - A_p : values greater than K_{p-1}



5. every non-terminal node – at least $\left\lceil \frac{m}{2} \right\rceil$ subtrees
- obs. limits on number of subtrees (and values) / node result from the manner in which inserts / deletes are performed so that the second requirement in the definition is met

* Example - B-tree of order 5

- non-terminal, non-root node – at most 5, at least 3 subtrees, i.e., between 2 and 4 values



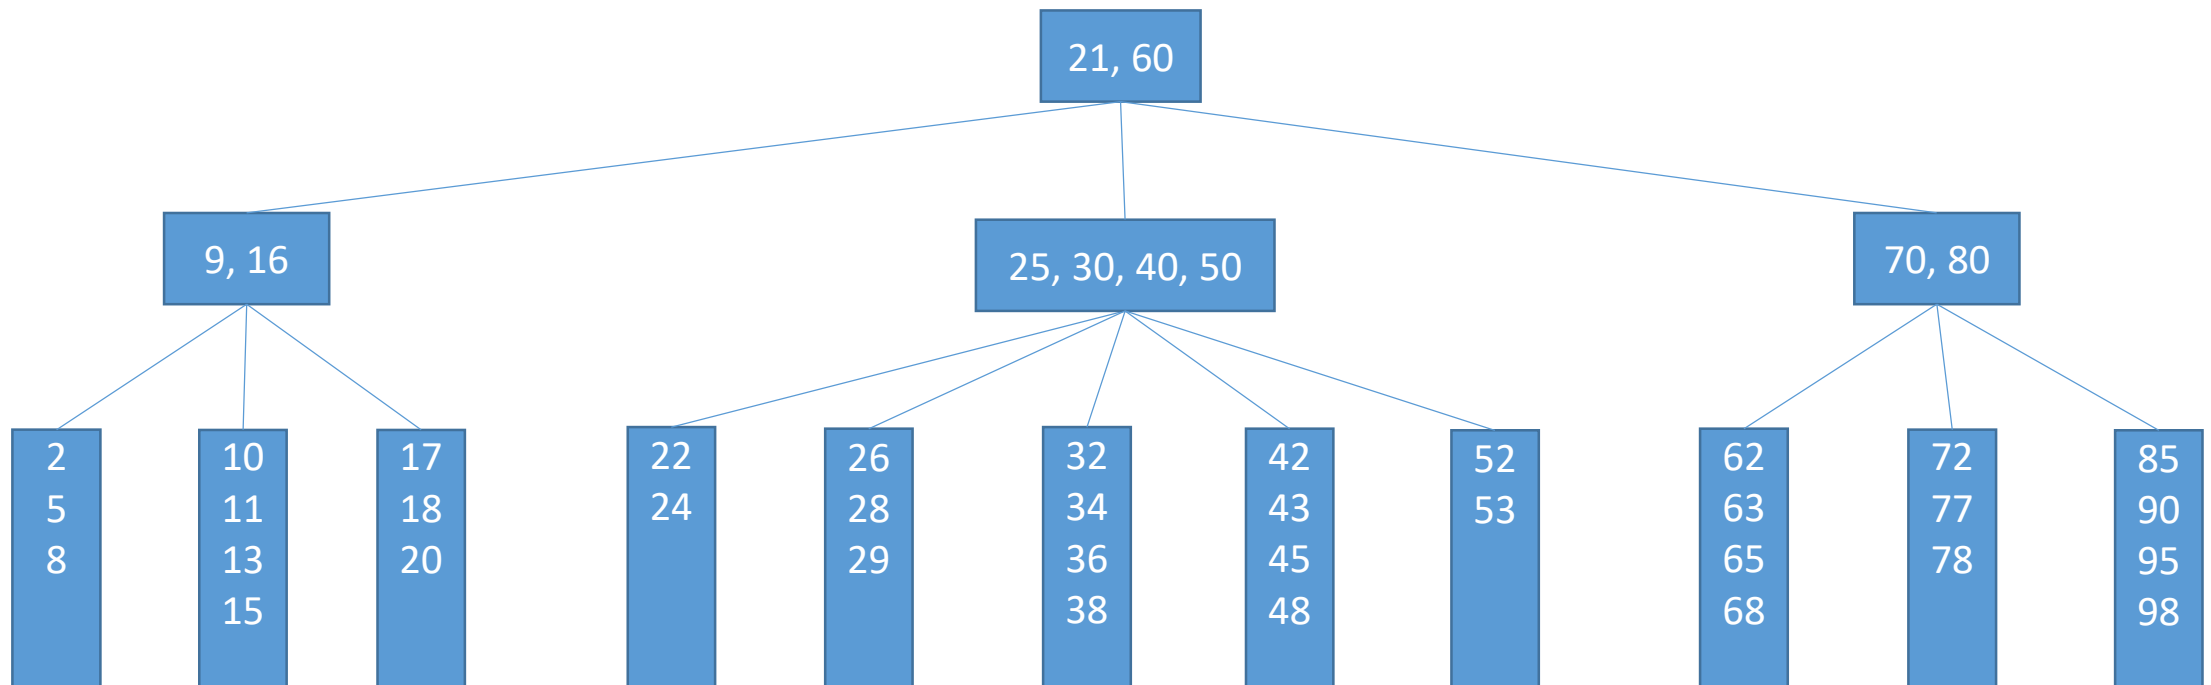
- B-tree of order m
 - storing the values of a key (a database index)
 - tree
 - key value + address of record
- 1. transformed into a binary tree
 - 2-3 tree method - generalization
- 2. the memory area allocated for a node can store the maximum number of values and subtree addresses

NV	K_1	$ADDR_1$...	K_{m-1}	$ADDR_{m-1}$	$Pointer_1$...	$Pointer_m$
----	-------	----------	-----	-----------	--------------	-------------	-----	-------------

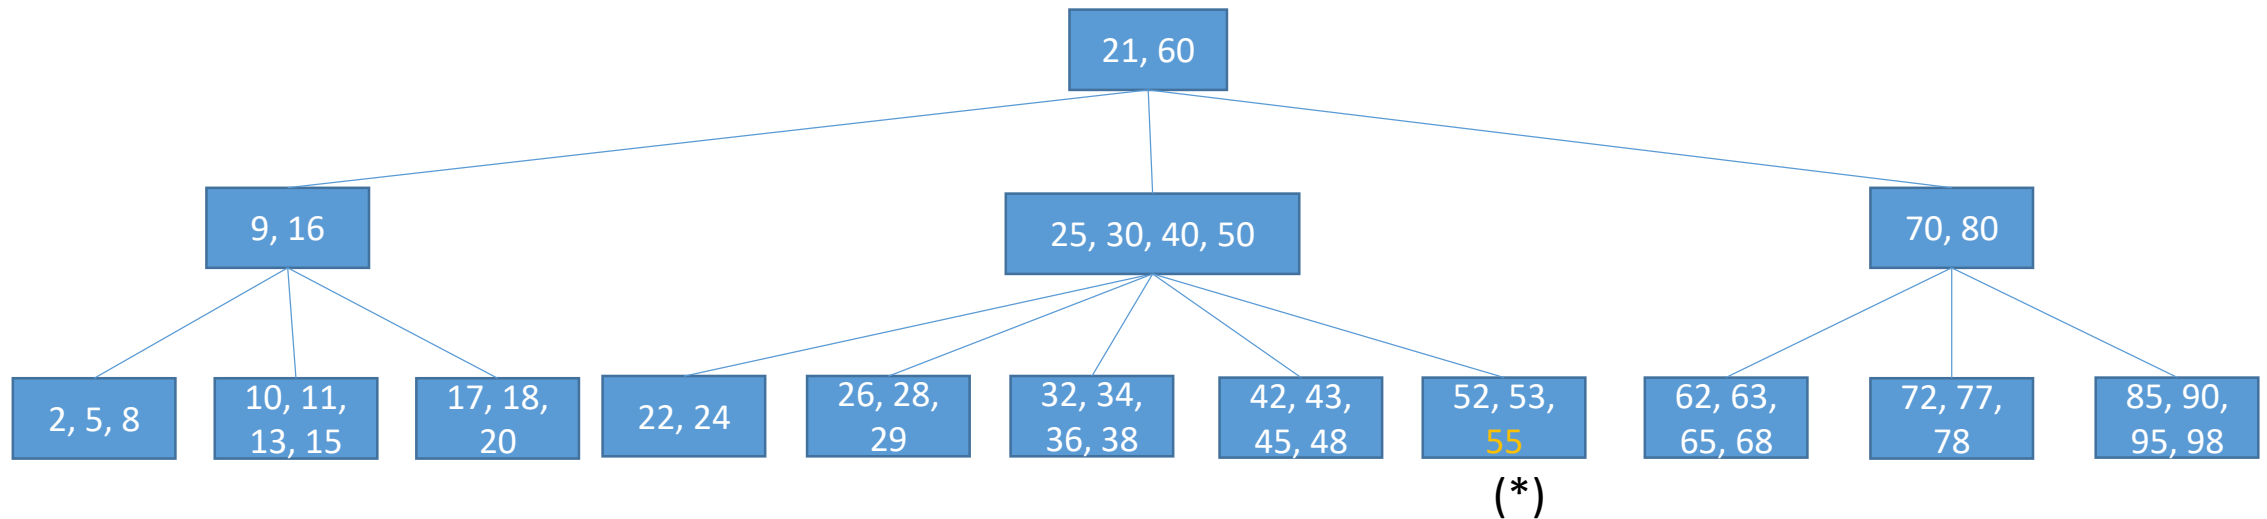
- NV - number of values in the node
- K_1, \dots, K_{m-1} - key values
- $ADDR_1, \dots, ADDR_{m-1}$ - the records' addresses (corresponding to the key's values)
- $Pointer_1, \dots, Pointer_m$ – subtree addresses

- B-tree of order m
 - useful operations in a B-tree
 - searching for a value
 - adding a value
 - removing a value
 - tree traversal (partial, total)

- B-tree of order m
 - adding a new value
 1. values in the tree must be distinct, i.e., the new value should not exist in the tree; perform a test, i.e., search for the value in the tree
 - if the new value can be added, the search ends in a terminal node
 2. if the reached terminal node has less than $m-1$ values, the new value can be stored in the node, e.g., 55 is added to the tree below:

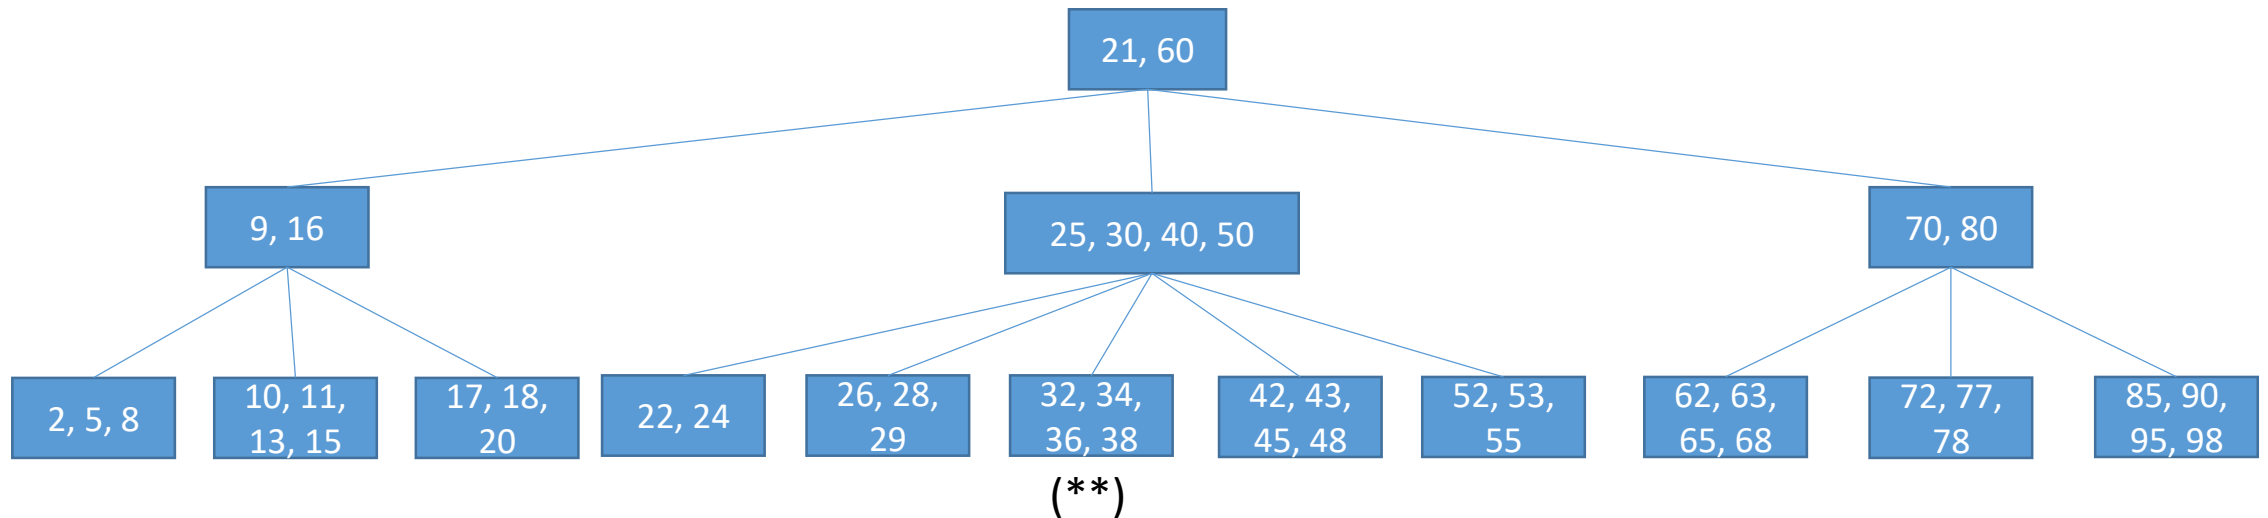


- B-tree of order m
 - adding a new value
 - the resulting tree is shown below:



- 55 belongs to the node marked with (*), which can store at most 4 values

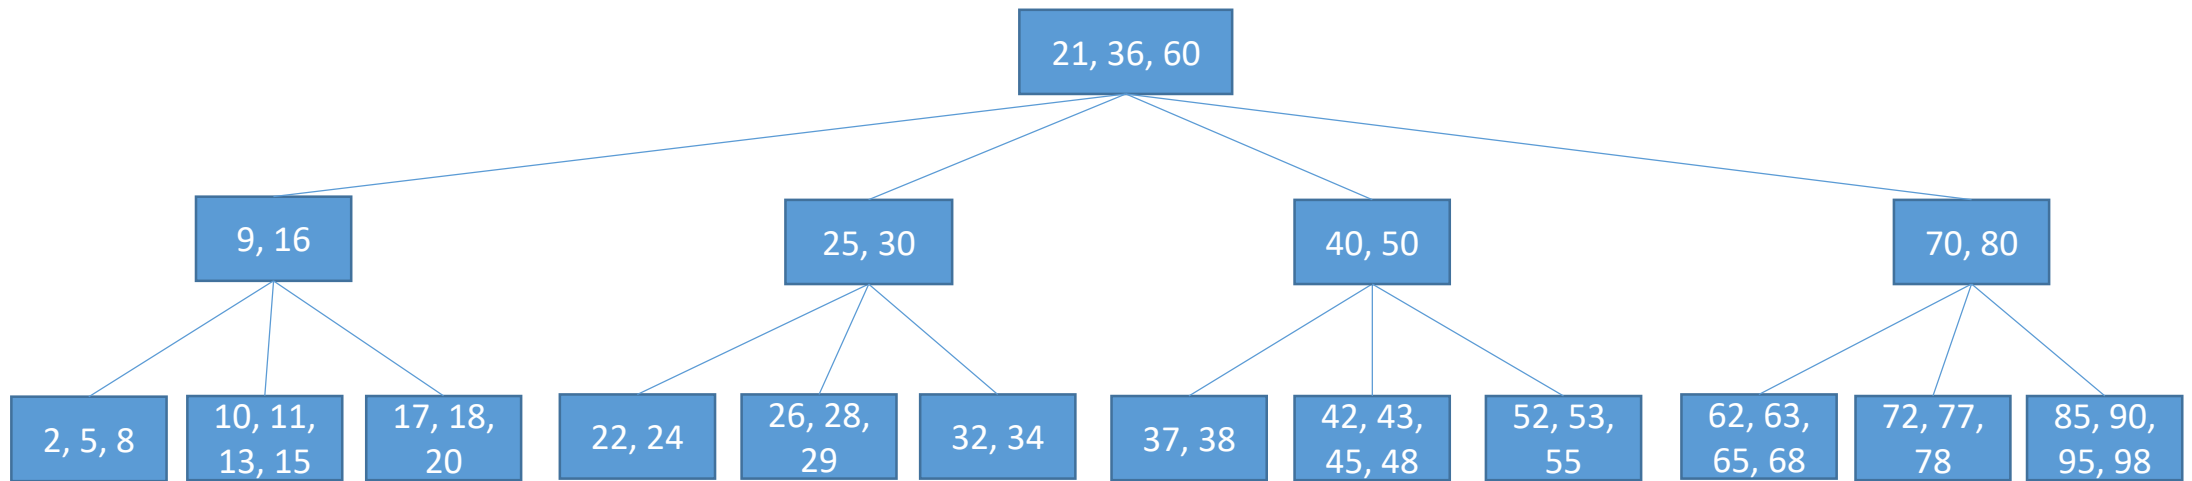
- B-tree of order m
 - adding a new value
 3. if the terminal node already has $m-1$ values, the new value is attached to the node, the m values are sorted, the node is split into 2 nodes, and the middle value (median) is attached to the parent node
 - e.g., add 37 to the tree below




- the node marked with (**) should contain values 32, 34, 36, 37, 38

->

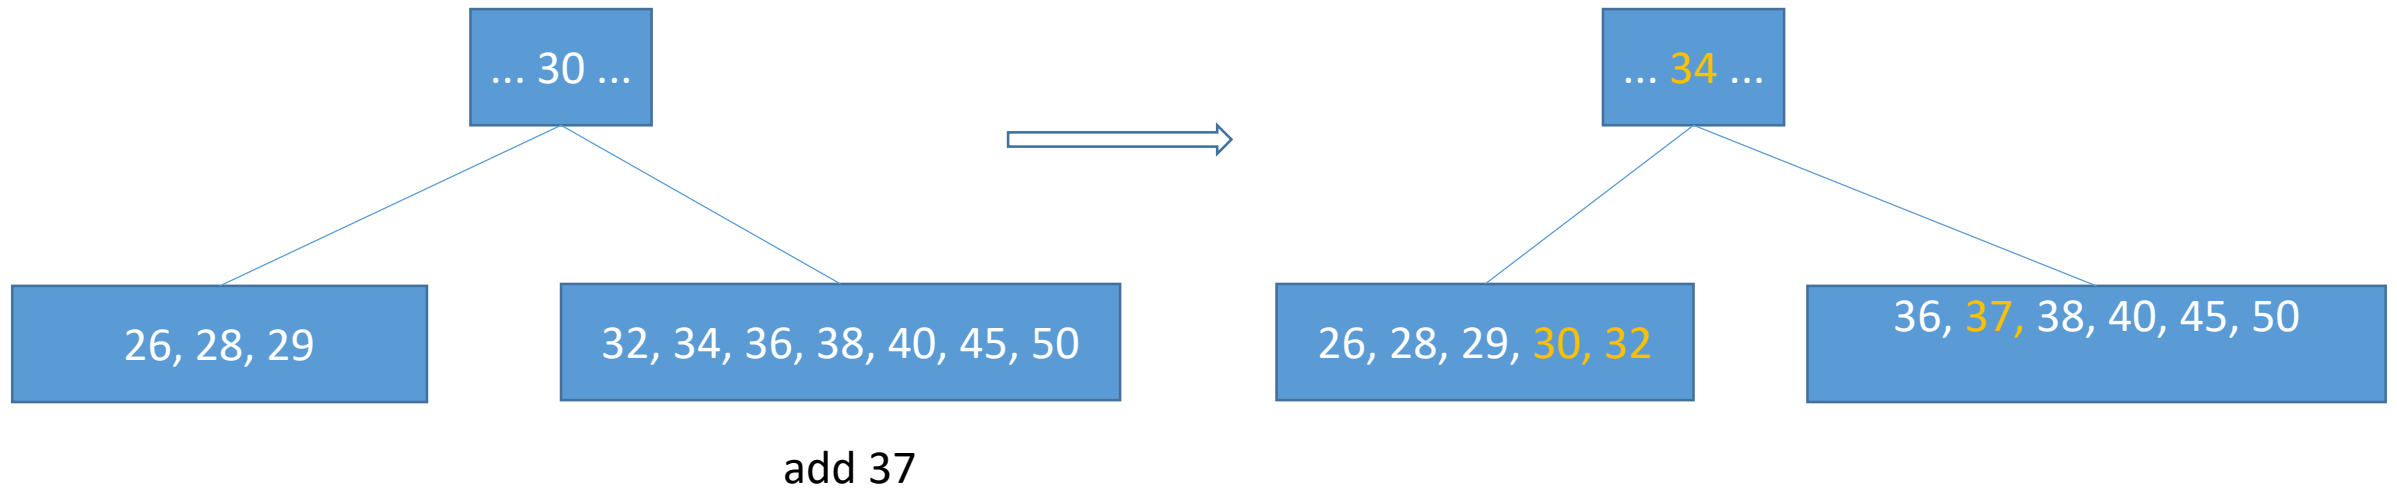
- B-tree of order m
 - adding a new value
 - since the node's capacity is exceeded, it is split into nodes 32, 34, and 37, 38, and 36 is attached to the parent node (with values 25, 30, 40, 50)
 - in turn, the parent must be split into 2 nodes (values 25, 30, and 40, 50), and 36 is attached to its parent



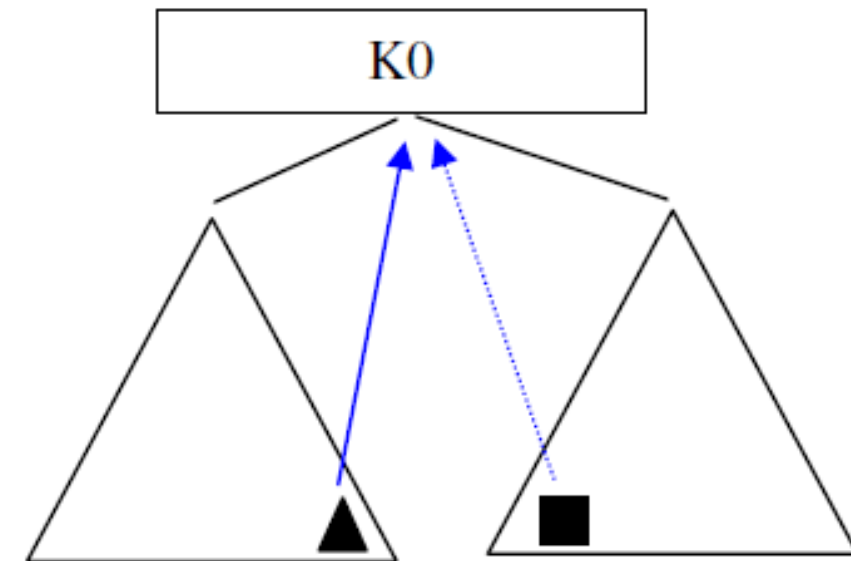
- B-tree of order m
 - adding a new value
 - optimizations
 - before performing a split - analyze whether one or more values can be transferred from the current node (with $m-1$ values) to a sibling node
 - e.g., B-tree of order 5 (non-terminal node - between 2 and 4 values, i.e., between 3 and 5 subtrees):



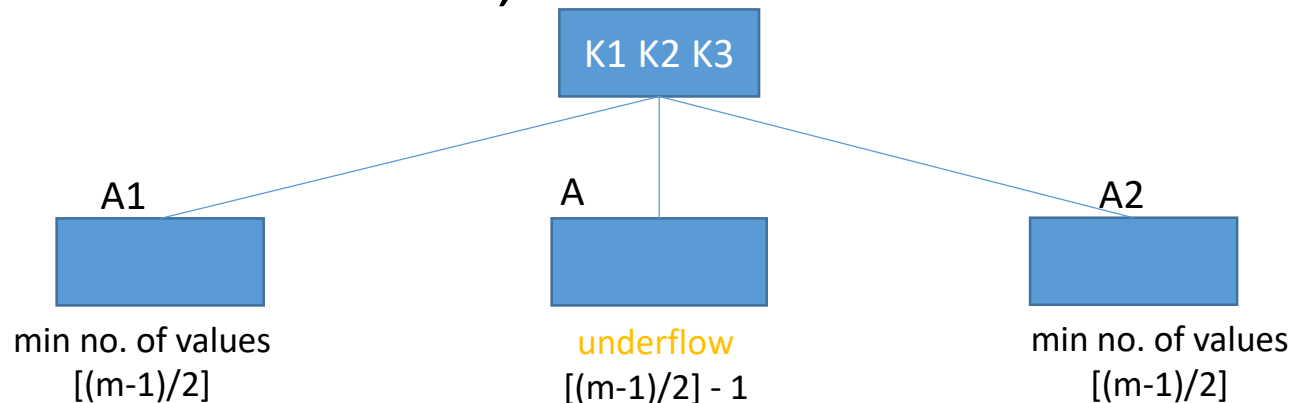
- B-tree of order m
 - adding a new value
 - optimizations
 - e.g., B-tree of order 8 (non-terminal node - between 3 and 7 values, i.e., between 4 and 8 subtrees):



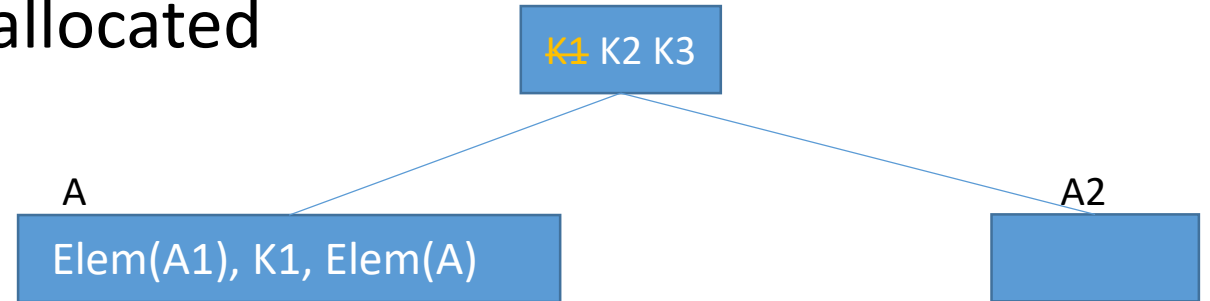
- B-tree of order m
 - removing a value
 - a node can have at most m subtrees, i.e., a maximum of m-1 values, and at least $\lceil \frac{m}{2} \rceil$ subtrees, i.e., at least $\lceil \frac{m}{2} \rceil - 1 = \lceil \frac{m-1}{2} \rceil$ values
 - when eliminating a value from a node, an underflow can occur (the node can end up with less values than the required minimum)
 - eliminate value K_0
 1. search for K_0 ; if it doesn't exist, the algorithm ends
 2. if K_0 is found in a non-terminal node (like in the figure on the right), K_0 is replaced with a *neighbor value* from a terminal node (this value can be chosen between 2 values from the trees separated by K_0)



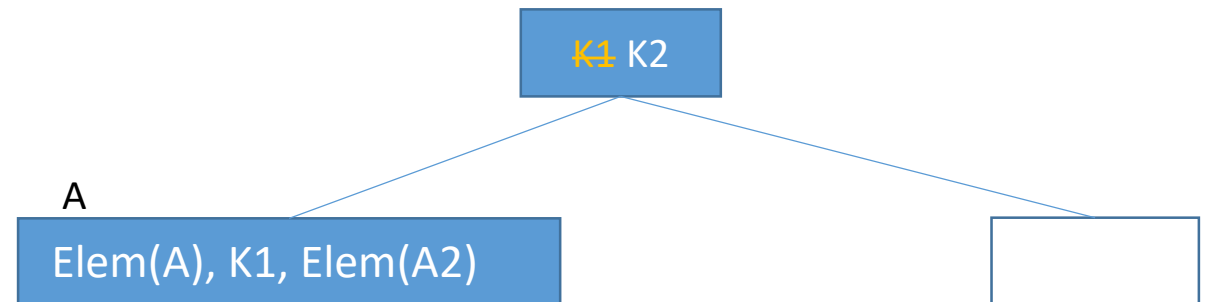
- B-tree of order m
 - removing a value
 3. perform this step until case a / b occurs
 - a. if the current node (from which a value is removed) is the root or underflow doesn't occur, the value is eliminated; the algorithm ends
 - b. if the delete operation causes an underflow in the current node (A), but one of the sibling nodes (left / right - B) has at least 1 extra value, values are transferred between A and B via the parent node; the algorithm ends
 - c. if there is an underflow in A, and sibling nodes A1 and A2 have the minimum number of values, nodes must be concatenated:



- B-tree of order m
 - removing a value
 - if A1 exists, A1 is merged with A and value K1 (separating A1 from A); the node at address A1 is deallocated



- if there is no A1 (A is the first subtree for its parent), A is merged with A2 and K1 (separating A from A2); the node at address A2 is deallocated



- case 3 is then analyzed for the parent node
- if the root is reached and has no values, it is removed and the current node becomes the root

- B-tree of order m
 - optimizations

obs 1. a block stores a node from a B-tree

- e.g.:
 - key size: 10b
 - record address / node address: 10b
 - NV value (number of values in the node): 2b
 - block size: 1024b (10b for the header)
- then: $2 + (m-1) * (10+10) + m * 10 = 1024 - 10 \Rightarrow m = 34$
- if the size of a block is 2048b and the other values are unchanged, then the order of the tree is $m = 68$, i.e., a node can have between 33 and 67 values

- B-tree of order m
 - optimizations
- the maximum number of required blocks (from the file that stores the B-tree) when searching for a value - the maximum number of levels in the tree; for $m=68$, if the number of values is 1.000.000, then:
 - the root node (on level 0) contains at least 1 value (2 subtrees)
 - on the next level (level 1) - at least 2 nodes * 33 values/node = 66 values
 - level 2 – at least 2^{*34} nodes * 33 values/node = 2.244 values
 - level 3 – at least 2^{*34*34} nodes * 33 values/node = 76.296 values
 - level 4 – at least $2^{*34*34*34}$ nodes * 33 values/node = 2.594.064 values, which is greater than the number of existing values => this level does not appear in the tree

=> at most 4 levels in the tree

- after at most 4 block reads and a number of comparisons in the main memory, it can be determined whether the value exists (the corresponding record's address can then be retrieved) or the search was unsuccessful

- B-tree of order m
 - optimizations


obs 2. in terminal nodes, subtree addresses are null; the space allocated for these addresses could be used to store additional (K, Addr) pairs

B+ tree

B+ tree

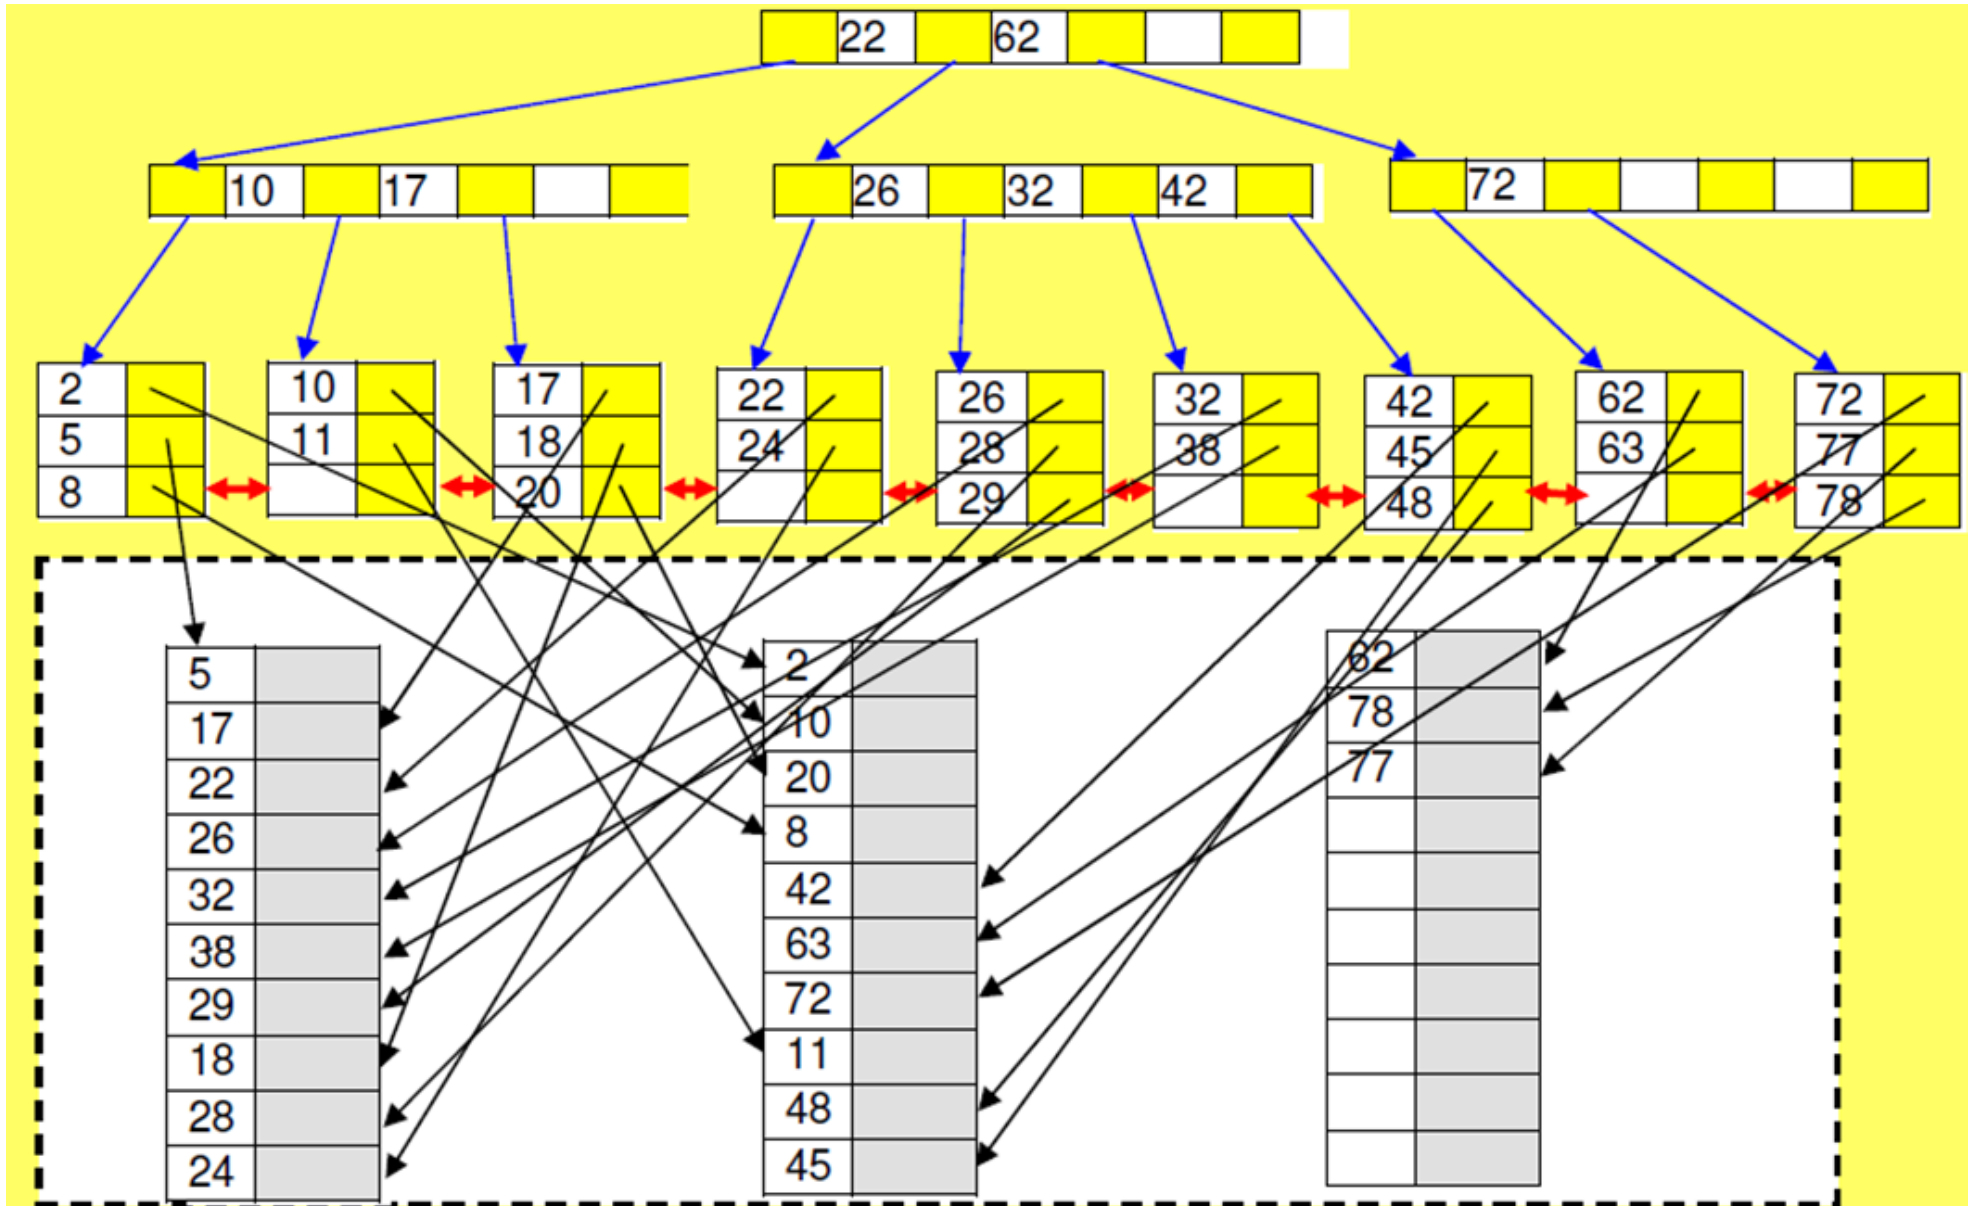
- B-tree variant
- the last level contains all the possible values (key values and the corresponding records' addresses)
- some key values can also appear in non-terminal nodes, without the records' addresses; their purpose is to separate values from terminal nodes (guide the search)
- terminal nodes are maintained in a doubly linked list; the data can be easily scanned via these links
- an inner node has the same min. / max. number of values as a B-tree
- a terminal node has at least $\lceil m / 2 \rceil$ values (instead of $\lceil (m - 1) / 2 \rceil$) and at most $(m-1)$ values
- for the min. / max. number of values, check obs 2. related to B-tree optimizations

B+ tree



- B+ tree of order 4

- B+ tree of order 4



B+ tree

- storing a B+ tree
 - B-tree methods
- operations (algorithms)
 - B-tree

B+ tree - in practice

* prefix key compression

- larger key size => less index entries fit on a page, i.e., less children / index page => larger B+ tree height
- keys in index entries - just direct the search => often, they can be compressed
- adjacent index entries with search key values: *Meteiut*, *Mircqkjt*, *Morqwkj*
- it's enough to store the values *Me*, *Mi*, etc
- what if the subtree also contains *Micfgjh*? => need to store *Mir* (instead of *Mi*)
- it's not enough to analyze neighbor index entries *Meteiut* and *Morqwkj*; the largest value in *Mircqkjt*'s left subtree and the smallest value in its right subtree also need to be examined
- inserts / deletes - modified correspondingly

B+ tree - in practice

- concept of *order* - relaxed, replaced by a physical space criterion, e.g., nodes should be at least half-full
- terminal / non-terminal nodes - different numbers of entries; usually, inner nodes can store more entries than terminal ones
- variable-length search key \Rightarrow variable-length entries \Rightarrow variable number of entries / page
- if a3 is used ($\langle k, \text{rid_list} \rangle$) \Rightarrow variable-length entries (in the presence of duplicates), even if attributes are of fixed length

B+ tree - in practice

- values found in practice
 - order – 200
 - fill factor (node) – 67%
 - fan-out – 133
 - capacity
 - height 4: $133^4 = 312,900,721$ records
 - height 3: $133^3 = 2,352,637$ records
- top levels can often be kept in the BP
 - 1st level – 1 page (8KB)
 - 2nd level – 133 pages (\cong 1MB)
 - 3rd level – 17689 pages (\cong 133 MB)

B+ tree - benefits

- balanced index => uniform search time
- rarely more than 3-5 levels, the top levels can be kept in main memory => only a few I/O operations are needed to search for a record
- widely used in DBMSs, among the most optimized components
- ideal for range selections, good for equality selections as well

References

- [Ra00] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (2nd Edition), McGraw-Hill, 2000
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007,
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Ta13] ȚÂMBULEA, L., Curs Baze de date, Facultatea de Matematică și Informatică, UBB, 2013-2014
- [Si10] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, McGraw-Hill, 2010
- [Ga08] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book, Prentice Hall Press, 2008